

Getting started

This document is broken into a few sections. First of all, there is information about getting a Python installation up and running. This is followed by a section on the basic usage of these programs to manipulate, plot, and execute simple non-negative least squares fits on GCMS datasets. The final section talks about generating calibration curves and automating data extraction with this calibration information. You can skip to that section if all you want to do is automate some data extractions. (It is not necessary to understand the data manipulation/plotting as that is automated in the final section.)

Installation

To get all of the files from this repository, you should be able to download a zip file of the current state of the repository. (Look for the button shown below on the [main page of this repository](#).) Unzip this repository wherever you like.




Figure 1: Download zipfile button

In addition, you will also need a Python 2.7 installation with the libraries Numpy (v 1.9.1 tested), matplotlib (v 1.4.2 tested), netCDF4 (v 1.0.4 tested), PyTables (v 3.1.1 tested), and Scipy (v 0.14.0 tested). Although not required, IPython (v 2.3.1 tested) provides a very useful advanced interactive Python interpreter, and examples in this documentation assume that you are using this environment.

Installing all of these packages can be time-consuming and error-prone, so it is best to start with an all-in-one package. In this case, the [Anaconda Python distribution](#) is a very nice option. It combines a large number of Python packages for scientific data analysis and a program (`conda`) for managing package updates (in addition to many other advanced features). The Anaconda developers (Continuum Analytics) have a lot of useful documentation for [installing Anaconda](#) and [using conda](#). On Windows, Anaconda may not install netCDF4. In this case, you can get a prebuilt installer from [Christoph Gohlke](#): be sure to get the Python 2.7 (“cp27”) 64-bit (“amd64”) build for the most recent version.

Learning the usage of all of these packages is far beyond the scope of this document. However, excellent documentation for most of the packages as well as full tutorials are [easily discovered](#).

Basics of working with GCMS data files

Export the data

First of all, be sure to export your GCMS data in a common data format, such as [AIA](#), [ANDI](#), or [CDF](#). It turns out that all of these formats [are related](#) in that they are all based off of [Network Common Data Format \(netCDF\)](#), so they may have the file extension “AIA” or “CDF”. This file type may not be the default for your instrument, so consult the documentation for your GCMS software to determine how to export your data in these formats.

Set up the processing environment

In order to process these files, run IPython from a terminal (command prompt in Windows) in the folder containing the “gcms.py” file (which is the base folder for this repository). There are (at least) two ways to do this that involve the command `cd` (change directory) run from either the terminal or an IPython session. For example (`home>$` is the command prompt located in the folder “home”, `In:` is the IPython prompt):

```
home>$ ipython
In: %cd "path-to-gcms-folder"
Out: path-to-gcms-folder
In:
```

or:

```
home>$ cd path-to-gcms-folder
gcms>$ ipython
In:
```

The “path-to-gcms-folder” is a valid path to the folder with “gcms.py”. It can take a little practice, but this gets easier very quickly. I’ll assume you use the second form of this as it makes using the [IPython notebook](#) much easier later.

Read AIA data files

First of all, you will need to `import` the “gcms.py” file to make the code accessible to the IPython environment. This file contains a class called `AIAFile` that reads and processes the GCMS files. `AIAFile` takes one argument, which is a string with the file name. This string must have the path (i.e. folder) information if the file is not in the same directory as “gcms.py”. Sample data files are contained in a folder called “data”.

```
In: import gcms
In: data = gcms.AIAFile('data/datasample1.CDF')
```

The variable `data` now contains our processed GCMS data set. You can see its contents using tab completion in IPython (`<tab>` refers to the tab key).

```
In: data.<tab>
data.filename data.intensity data.nnls data.ref_build data.times
data.integrate data.masses data.tic
```

All of these attributes are either data that describe or functions that modify (methods) our dataset. You can inspect these attributes very easily in IPython by just typing the name at the prompt.

```
In: data.times
Out:
array([0.08786667, ..., 49.8351])
In: data.tic
Out:
array([158521., ..., 0.]
```

This is a short description of these initial attributes:

- *filename*: This is the name of the file that you imported.
- *times*: A Numpy array of the times that each MS was collected.
- *tic*: A Numpy array of the total ion chromatogram intensities.
- *masses*: A Numpy array the masses that cover the data collected by the MS.
- *intensity*: This is the 2D Numpy array of raw MS intensity data. The columns correspond to the masses in the `masses` array and the rows correspond to the times in the `times` array.

The remaining attributes `ref_build`, `nnls`, and `integrate` are functions that deal with the non-negative fitting routine and are covered in later sections.

Simple plotting

We can easily plot these data using the plotting package Matplotlib. As an example, let's try making a plot of our total ion chromatogram. In this case, `data.times` will be our "x-axis" data, and `data.tic` will be our "y-axis" data.

```

In: import matplotlib.pyplot as plt
In: plt.plot(data.times, data.tic)
Out:
[<matplotlib.lines.Line2D at 0x7f34>]
In: plt.show()

```

This should produce a pop-up window with an interactive plot of your TIC. (This process should be fairly quick. However, sometimes the plot initially appears behind the other windows, which makes it seem like things are stuck. Be sure to scroll through your windows to find it.)




Figure 2: Total ion chromatogram

One drawback here is that you have to type these commands every time you want to see this plot. Of course, you can simply put all of these commands into a text file and run it with Python directly. Copy the following code into a file called “tic_plot.py”.

```
import matplotlib.pyplot as plt
```

```
import gcms

data = gcms.AIAFile('data/datasample1.CDF')
plt.plot(data.times, data.tic)
plt.show()
```

It is common practice to do all imports at the top of a Python program. That way it is clear exactly what code is being brought into play. Run this new file using the `python` command from the terminal.

```
gcms>$ python tic_plot.py
```

The window with your plot will now appear. (You will not be able to work in the terminal until you close this window.) Alternatively, you can run this program directly from IPython.

```
gcms>$ ipython
In: %run tic_plot.py
```

This also pops open a new window containing the interactive plot. It has the advantage, however, that once the window is closed, you are dropped back into an IPython session that “remembers” all of the variables and imports that you created in your program file. In our example above, once the plot window is closed, your IPython session will have `gcms`, `plt`, and `data` (our GCMS AIA file) available. This is very useful if you want to continue to work interactively with your data, and it is a great way to remove a bunch of repetitive typing.

Working with multiple data sets

In our example above, we opened our dataset into a variable called `data` in order to be able to plot the TIC. If you want to manipulate more than one data set, the procedure is exactly the same, except that you will need to use different variable names for your other data sets.

```
In: data2 = gcms.AIAFile('data/datasample2.CDF')
```

These two data sets can be plot together on the same figure by doing the following:

```
In: plt.plot(data.times, data.tic)
Out:
[<matplotlib.lines.Line2D at 0x7f34>]
In: plt.plot(data2.times, data2.tic)
```

```
Out:
[<matplotlib.lines.Line2D at 0x02e3>]
In: plt.show()
```

The following window should appear on the screen. (There is a blue and green line here that are a little hard to see in this picture. Zoom in on the plot to see the differences.)

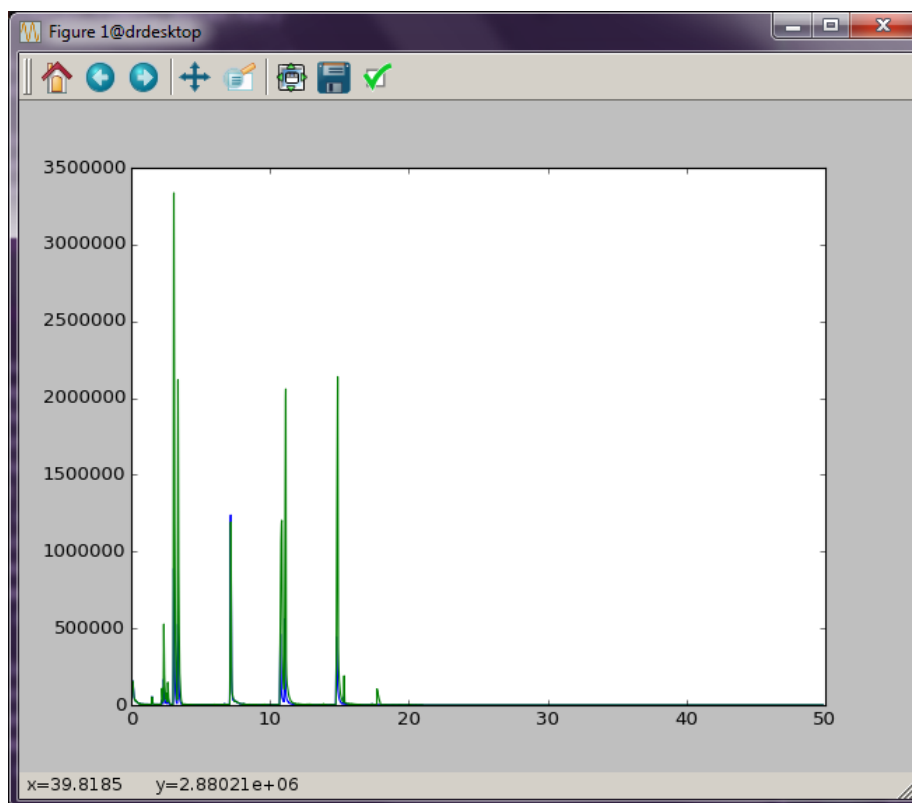


Figure 3: Two tic plotted together

Reference MS Files

First of all, you are going to need a series of reference MS spectra to do the fitting. I got my references from [massBank](#). The format of these files is important, so a couple of examples are provided in this repository (refcpd.txt and refcpd2.txt). The most important aspect of these file is that there are two space-separated data points per line. The first data point is the mass and the second data point

is an intensity. (The intensities are normalized in the fitting program, so it is not necessary to make those adjustments yourself.) Comments can be added by starting a line with '#'. This must be the first character; don't put white space before this symbol. Use comment lines to keep track of where the MS data was acquired and any identifying data.

To use these files for fitting, they need to be added to the file 'reference_files.txt'. Again, an example of this file is provided. The ordering of compounds in this file is important. Any time you make changes to this file, you must rerun both calibration.py and data.py to regenerate the data files that are created by these programs.

Calibration Data

If you have calibration data for a particular reference compound, you must create a csv file and folder that have the same base name as the reference MS file from above. Again, an examples are provided in this repository called refcpd.csv and the folder refcpd. All of your calibration AIA files for this compound need to be stored in the newly created folder. In order for these new data files to be processed, the refcpd.csv file must be appropriately modified.

The csv file is a simple comma-separated text file, but again the structure is important. The first row in this file is critical. At the end, there are two values that define the starting and stopping time points for integration. Change these values based on the time range that you've determined from the TIC of a calibration run. The rest of the rows are data file information. The first column is the name of a calibration data file, and the second column needs to be the concentration of the reference compound associated with that run. You don't have to add all of the calibration files here, but if they are not in this list, they won't be processed. Alternatively, any line that starts with a '#' is a comment, and will be ignored. In this way, you can comment out samples, and add some notes as to why that sample was not used or whatever.

Run Calibrations

Once you've updated the calibration information from above. You can run the program 'calibration.py'. This runs through all of the reference spectra defined in the 'reference_files.txt' file. If a 'csv' file exists for a particular reference file, then a calibration will be performed.

All of the calibration data files listed in the csv file will be processed and a calibration curve generated. For each calibration sample, a plot of the reference-extracted data will be generated in the calibration folder (refcpd_fits.png). In addition, a calibration curve plot is also generated (refcpd_cal_curve.png),

which plots the integrated intensities and calibrated intensities vs the concentrations. In addition, the calibration information is printed on the graph for quick visual inspection. There is no need to write down this calibration information.

This program has some important command line arguments that will change the programs defaults. The first argument, ‘-nobkg’, is a simple flag for background fitting. By default, the fitting routine will select a MS slice from the data set and use that as a background in the non-negative least squares fitting. This procedure can change the integrated values. If you use this flag, then a background MS will not be used in the fitting. Using a background slice in the fitting may or may not give good results. It might be a good idea to look at your data with and without the background subtraction to see which is better.

The second command line argument is ‘-bkg_time’. By default, the fitting program uses the first MS slice as a background for fitting. However, if there is another time that looks like it might make a better background for subtraction, then you can put that number here.

Here’s a couple of example usages of this script:

```
# This will run the calibration program with all defaults
$ python calibration.py
# This shuts off the background subtraction
$ python calibration.py --nobkg
# This sets an alternate time for the background subtraction
# In this case, the time is set to 0.12 minutes
$ python calibration.py --bkg_time 0.12
```

Another file is also generated during this process: cal.h5. This is a HDF5 file that contains all of the calibration information for each standard. Do not delete this file; it is essential for the next step. This is a very simple file, and there are many tools for looking at the internals of an HDF5 file. For example, [ViTables](#) is recommended. The background information, such as whether a background was used and the time point to use as a background spectrum, are stored as user attributes of the calibration table.

Process Sample Data

Put all of your data files in a folder that must be called ‘data’. Once you’ve done this, run the program ‘data.py’ to process every AIA data file in that folder using the calibration information that was determined from the steps above.

This program opens the AIA file for the sample and performs non-negative least squares analysis of the full data set using the reference spectra that are listed in the ‘reference_files.txt’ file. Using the calibration information that was determined above, it finds the concentrations of those components in the sample

data. For every reference compound that has associated calibration information, a plot is generated that overlays the TIC (gray) and extracted reference fit (blue). The title of the plot provides the calibrated concentration information. Visual inspection of these files is recommended.

This file also accepts the same command line arguments as ‘calibration.py’ from the section above. You will be warned if you try to analyze your data with different background information than the calibration samples. This may not impact your data much, but it is good to know if you are doing something different.

This file also generates another HDF5 file called ‘data.h5’, which contains the integration and concentration information for every component. This information is identical to what is printed on the extraction plots above. However, this tabular form of the data is a bit more convenient for comparing many data sets. See the Calibration section for a recommended HDF5 file viewer.